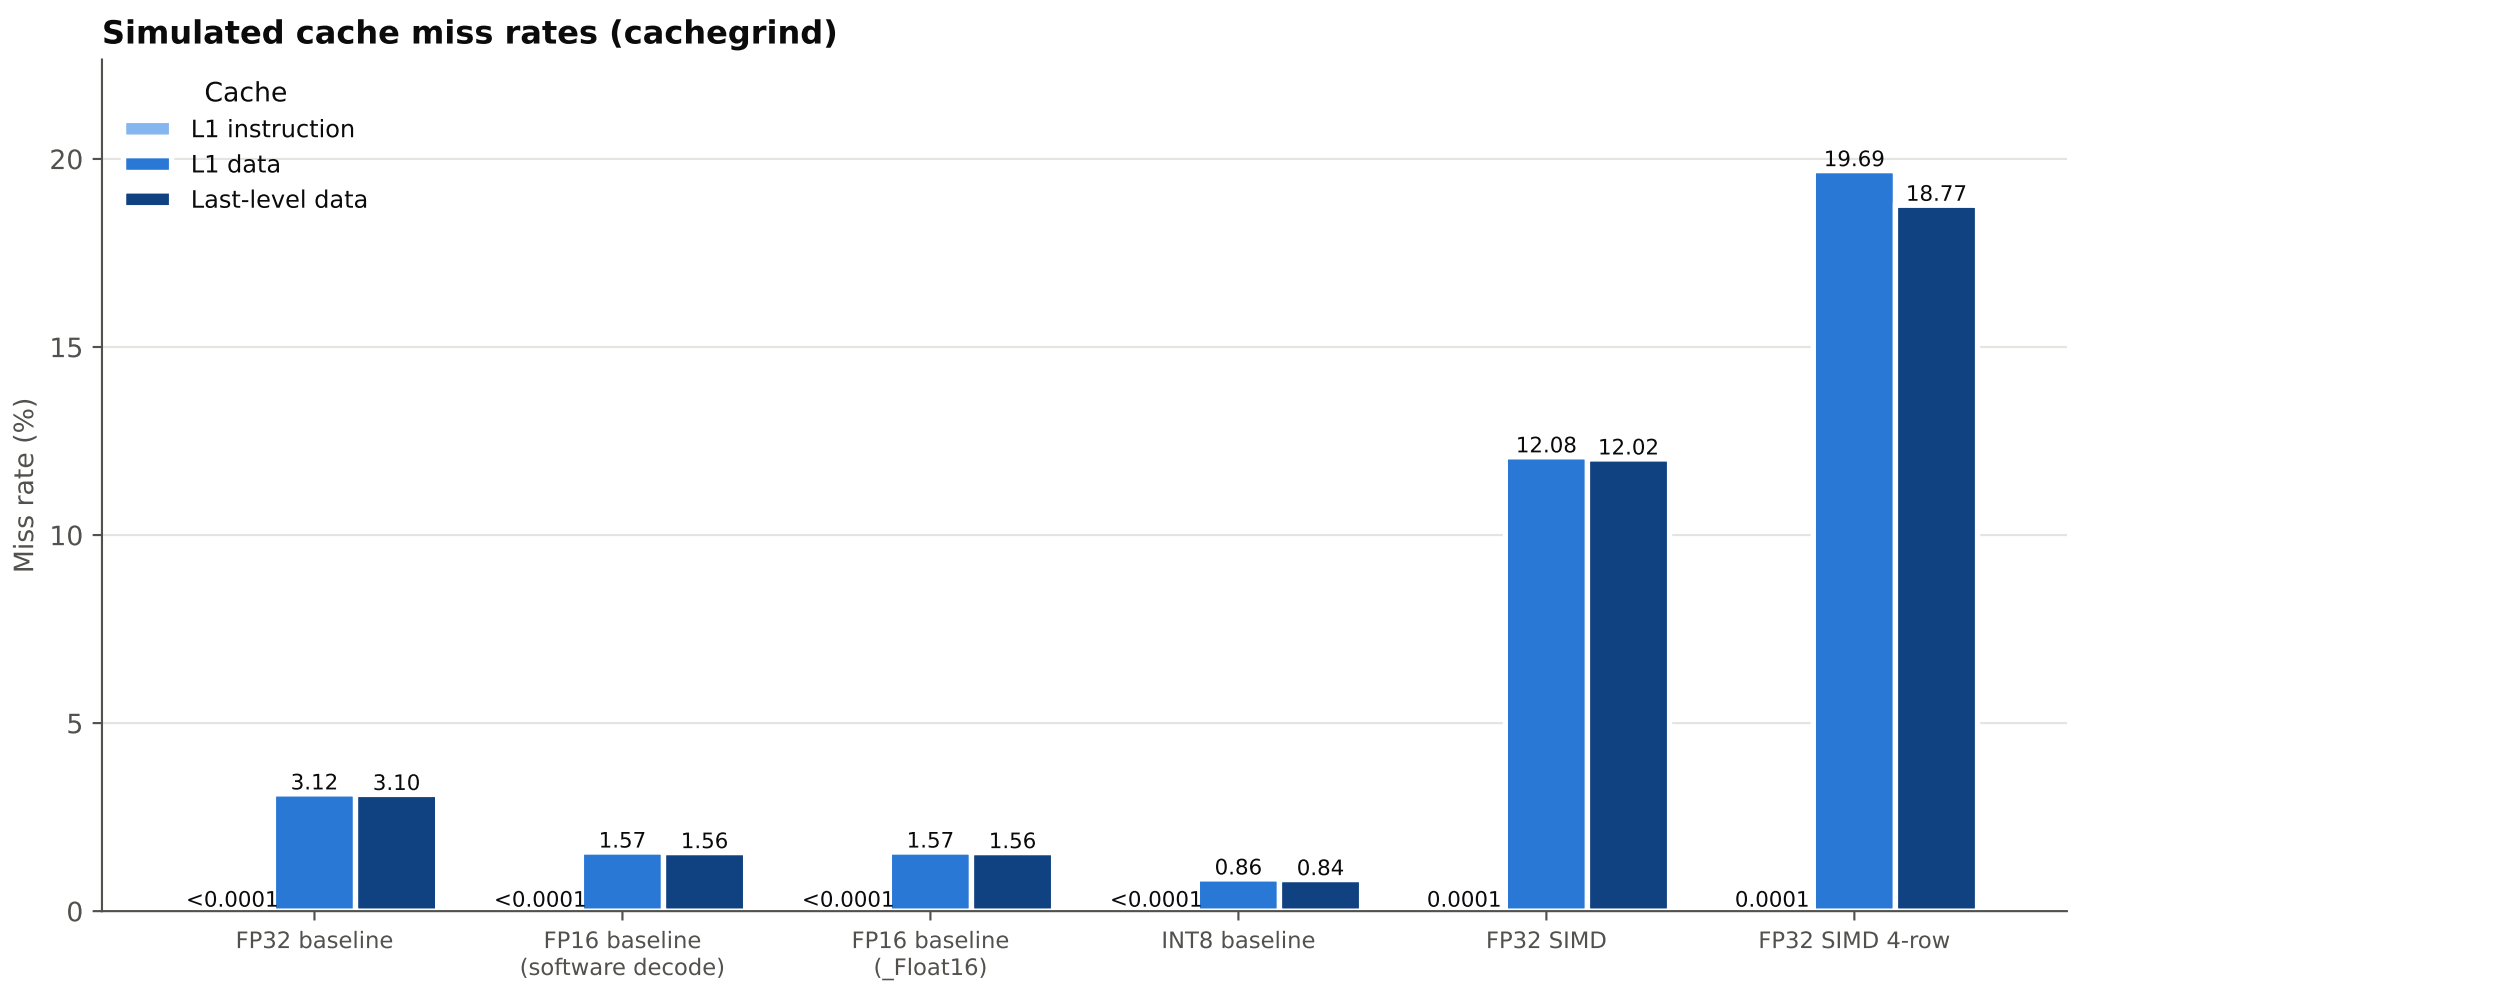

Cachegrind with a fixed 32 KiB L1 I/D and 8 MiB last-level cache, 16 tokens, 1 thread (OpenMP builds run with OMP_NUM_THREADS=1: valgrind cannot model per-core caches). Data miss rate = (read + write misses) / (reads + writes).