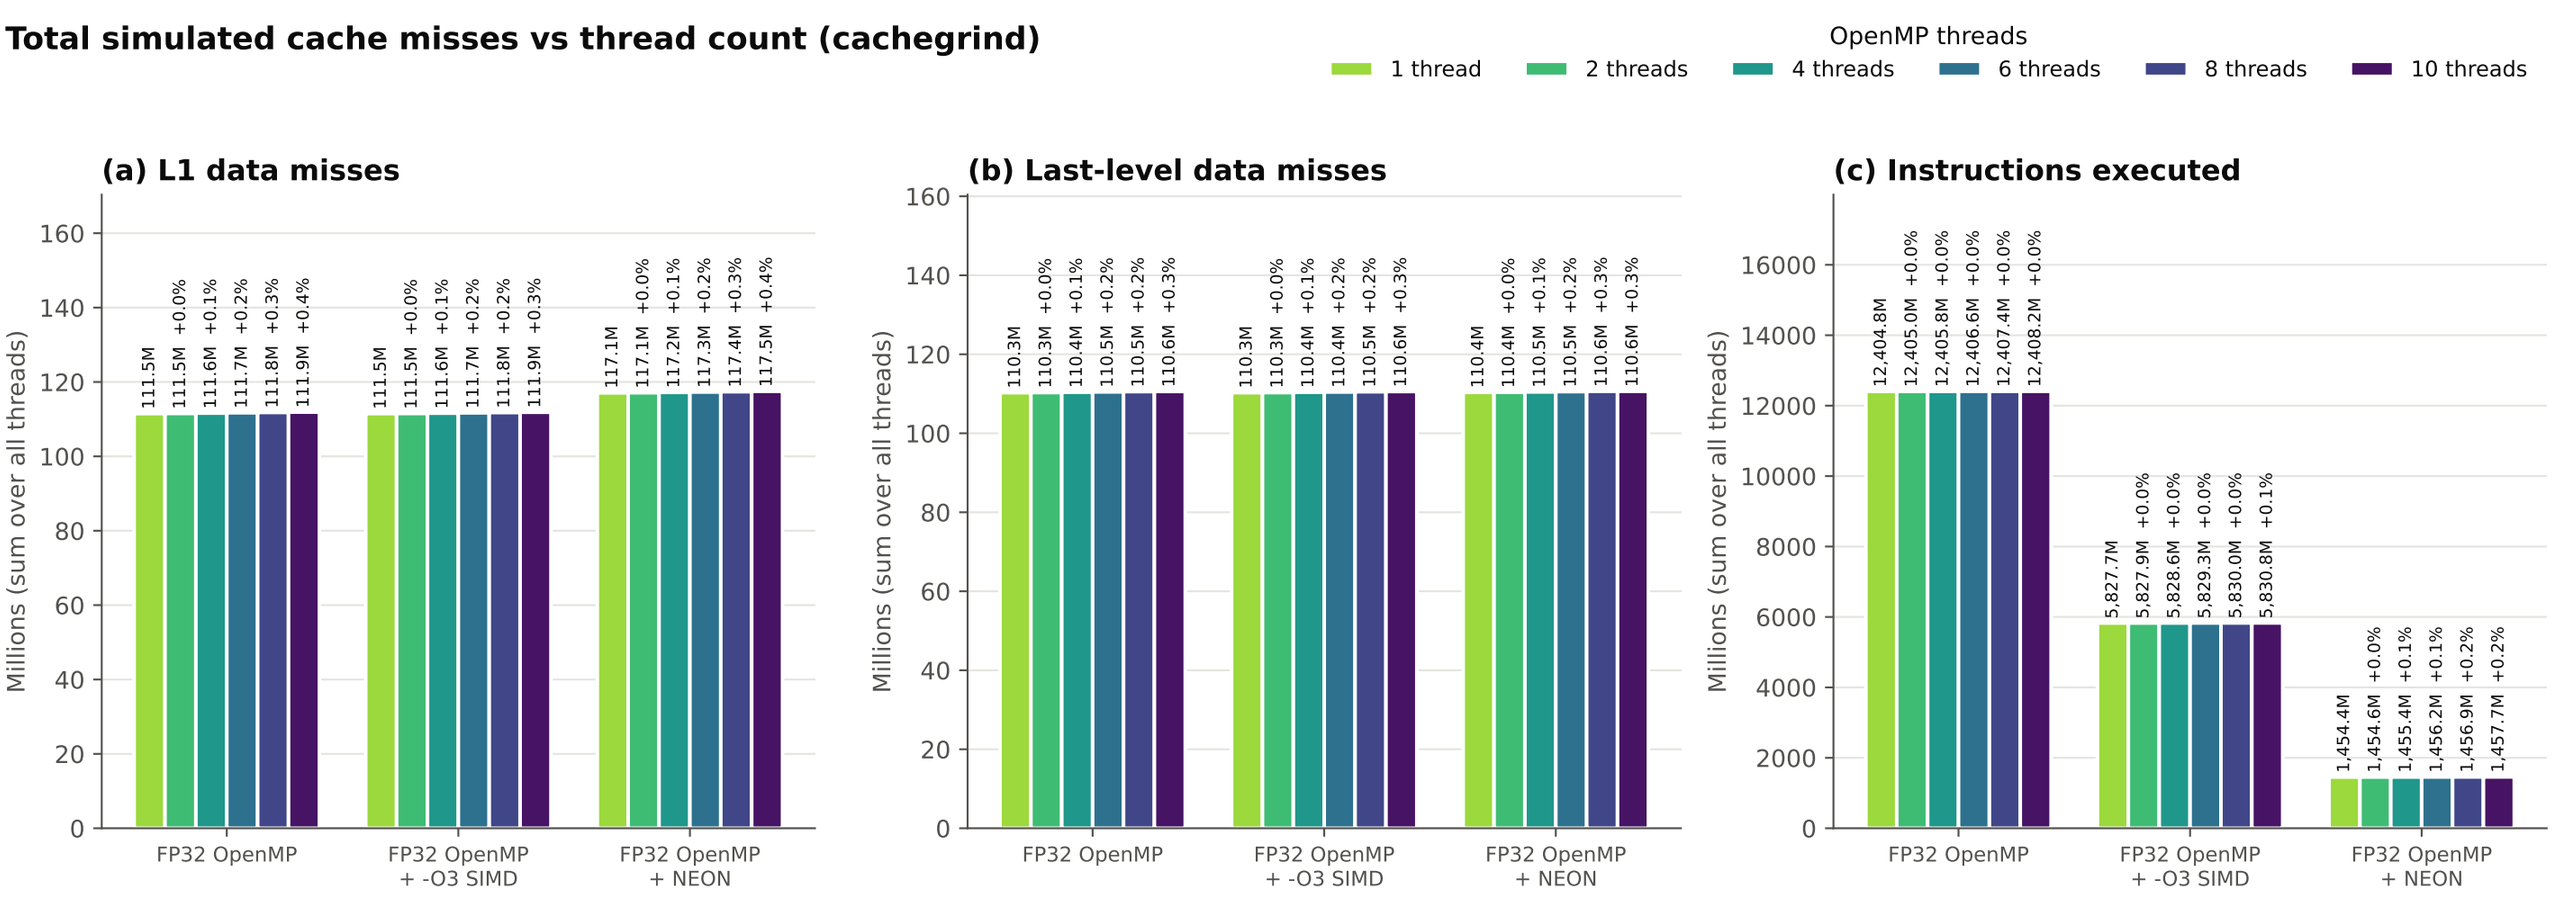

Sum over all threads, 16 tokens, OMP_WAIT_POLICY=passive. Valgrind runs threads one at a time on ONE simulated 32 KiB L1 / 8 MiB LL cache, so totals are comparable, per-core behaviour is not. % = change vs the same build at 1 thread.